

Group Assignment 3

February 22, 2021

0.1 Tracking Water Bill Debt by Demographics across California

0.1.1 Load Libraries and Data

First, I will load the libraries

```
[1]: import pandas as pd
import plotly.express as px
```

Now, I am going to upload the water bill data

```
[2]: acs = pd.read_csv('Data/Updated Bill Data 2_22.csv')
```

0.1.2 Explore the Data

Now that my data are uploaded, I need to get a sense of how they look.

```
[3]: acs.shape
```

```
[3]: (1073, 52)
```

This dataset has 1073 entries and 52 columns - it's big!

I also want to get a sense of the data itself: missing data, data type, etc.

```
[4]: acs.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1073 entries, 0 to 1072
Data columns (total 52 columns):
 #   Column                                Non-Null Count
Dtype
---  ---
-----
 0   Zip Codes                            1073 non-null
object
 1   Count of Zip Code                    1072 non-null
float64
 2   Sum of Less than $100                 934 non-null
float64
 3   Sum of $100-$200                      965 non-null
```

float64		
4	Sum of \$200-\$300	965 non-null
float64		
5	Sum of \$300-\$400	954 non-null
float64		
6	Sum of \$400-\$500	941 non-null
float64		
7	Sum of \$500-\$600	917 non-null
float64		
8	Sum of \$600-\$700	914 non-null
float64		
9	Sum of \$700-\$800	906 non-null
float64		
10	Sum of \$800-\$900	895 non-null
float64		
11	Sum of \$900-\$1000	894 non-null
float64		
12	Sum of More than \$1000	936 non-null
float64		
13	Sum of Total number of delinquent residential accounts	1063 non-null
float64		
14	pop	1030 non-null
float64		
15	nhw	1030 non-null
float64		
16	black	1030 non-null
float64		
17	hisp	1030 non-null
float64		
18	asian	1030 non-null
float64		
19	noncitizen	1030 non-null
float64		
20	immigrants	1030 non-null
float64		
21	ohu	1030 non-null
float64		
22	lep_hh	1030 non-null
float64		
23	dpov	1030 non-null
float64		
24	npov	1030 non-null
float64		
25	mhhi	1030 non-null
float64		
26	overcrowded	1030 non-null
float64		
27	no_veh_hh	1030 non-null

float64		
28	w_broadband	1030 non-null
float64		
29	pop_19_64	1030 non-null
float64		
30	uninsured_19_64	1030 non-null
float64		
31	pct_nhw	1030 non-null
float64		
32	pct_black	1030 non-null
float64		
33	pct_hisp	1030 non-null
float64		
34	pct_asian	1030 non-null
float64		
35	pct_noncitizen	1030 non-null
float64		
36	pct_immigrants	1030 non-null
float64		
37	pct_lep_hh	1029 non-null
float64		
38	pct_povt	1029 non-null
float64		
39	pct_overcrowded	1029 non-null
float64		
40	pct_no_veh_hh	1029 non-null
float64		
41	pct_broadband	1029 non-null
float64		
42	pct_no_broadband	1029 non-null
float64		
43	pct_uninsured_19_64	1030 non-null
float64		
44	aggveh	1030 non-null
float64		
45	pct_no_hins	1030 non-null
float64		
46	veh_person	1030 non-null
float64		
47	Total Population in Occupied Housing Units: Renter Occupied	1030 non-null
float64		
48	Owner Occupied Pop	1030 non-null
float64		
49	% Renter Pop	1030 non-null
float64		
50	% Owner Pop	1030 non-null
float64		
51	Households	1031 non-null

```
float64
dtypes: float64(51), object(1)
memory usage: 436.0+ KB
```

So here we can see that there are fewer datapoints here than total zip codes in California. There are over 1,700 zips in the state. This can largely be attributed to the fact that the water bill debt data was conducted via a survey distributed by the California State Water Resources Control Board. Survey responses have their limitations in that they are completed on a voluntary basis. It is worth keeping in mind as I continue with my analysis that this is not a complete dataset of all zip codes in the state. Hopefully, this dataset is complete enough though to draw some conclusions about water bill debt trends and demographics.

Note that all the data types are floats as well - so I shouldn't have problems conducting quantitative analyses.

0.1.3 Cleaning the Data

Next, I will only keep the columns of interest, including debt-related columns, as well as racial/ethnic factors, percent poverty, percent renter, percent immigrant, and median household income. I have a sense that these factors will contribute to making an area more financially vulnerable, and thus, might highlight trends in water bill debt data.

```
[5]: refined_columns = ['Zip Codes',
    'Sum of Less than $100',
    'Sum of $100-$200',
    'Sum of $200-$300',
    'Sum of $300-$400',
    'Sum of $400-$500',
    'Sum of $500-$600',
    'Sum of $600-$700',
    'Sum of $700-$800',
    'Sum of $800-$900',
    'Sum of $900-$1000',
    'Sum of More than $1000',
    'Sum of Total number of delinquent residential accounts',
    'pop',
    'mhhi',
    'pct_black',
    'pct_hisp',
    'pct_asian',
    'pct_povt',
    'pct_immigrants',
    '% Renter Pop']
```

```
[6]: acs = acs[refined_columns]
```

Now I have just saved the new data frame and will check my work.

```
[7]: acs.head()
```

```

[7]: Zip Codes Sum of Less than $100 Sum of $100-$200 Sum of $200-$300 \
0 90001 5726.0 4937.0 1582.0
1 90002 3130.0 2152.0 1052.0
2 90003 1829.0 1880.0 1562.0
3 90004 2199.0 1659.0 1119.0
4 90005 1712.0 1107.0 598.0

Sum of $300-$400 Sum of $400-$500 Sum of $500-$600 Sum of $600-$700 \
0 684.0 395.0 283.0 220.0
1 708.0 493.0 385.0 343.0
2 1169.0 941.0 782.0 673.0
3 735.0 502.0 387.0 279.0
4 401.0 281.0 175.0 146.0

Sum of $700-$800 Sum of $800-$900 ... Sum of More than $1000 \
0 137.0 132.0 ... 709.0
1 281.0 231.0 ... 1779.0
2 513.0 482.0 ... 3437.0
3 223.0 206.0 ... 1116.0
4 95.0 95.0 ... 426.0

Sum of Total number of delinquent residential accounts pop mhhi \
0 14902.0 58975.0 38521.0
1 10755.0 53111.0 35410.0
2 13669.0 72741.0 37226.0
3 8589.0 61586.0 48754.0
4 5094.0 39479.0 35149.0

pct_black pct_hisp pct_asian pct_povt pct_immigrants % Renter Pop
0 0.088648 0.900144 0.002187 0.287524 0.407630 0.647257
1 0.194950 0.784640 0.006063 0.328603 0.350737 0.621999
2 0.221828 0.770542 0.003533 0.306597 0.375442 0.692044
3 0.038856 0.511139 0.251437 0.180601 0.495437 0.803900
4 0.061374 0.492338 0.350794 0.280593 0.591352 0.912764

```

[5 rows x 21 columns]

Just the cell columns of interest have saved.

0.1.4 Normalize the Data

I need the percentage of the delinquent population each debt “bucket” represents.

I also want to know what percent of the population in each zip code has water bill debt, broadly. To do this, I need to create new columns for each of these columns as a percentage of the delinquent population and total population, respectively.

Further I want to convert all of the % values from decimals to percents so they are clearer on my maps.

```
[8]: list(acs)
```

```
[8]: ['Zip Codes',  
      'Sum of Less than $100',  
      'Sum of $100-$200',  
      'Sum of $200-$300',  
      'Sum of $300-$400',  
      'Sum of $400-$500',  
      'Sum of $500-$600',  
      'Sum of $600-$700',  
      'Sum of $700-$800',  
      'Sum of $800-$900',  
      'Sum of $900-$1000',  
      'Sum of More than $1000',  
      'Sum of Total number of delinquent residential accounts',  
      'pop',  
      'mghi',  
      'pct_black',  
      'pct_hisp',  
      'pct_asian',  
      'pct_povt',  
      'pct_immigrants',  
      '% Renter Pop']
```

```
[9]: acs['Percent Delinquent'] = acs['Sum of Total number of delinquent residential_  
    ↪accounts']/acs['pop']*100  
  
pct_debt_buckets = ['Percent Less than $100', 'Percent $100-$200', 'Percent_  
    ↪$200-$300' ,  
                    'Percent $300-$400', 'Percent $400-$500', 'Percent $500-$600',  
    ↪'Percent $600-$700',  
                    'Percent $700-$800', 'Percent $800-$900', 'Percent $900-$1000',  
    ↪'Percent More than $1000']  
  
debt_buckets = ['Sum of Less than $100', 'Sum of $100-$200', 'Sum of_  
    ↪$200-$300' ,  
                'Sum of $300-$400', 'Sum of $400-$500', 'Sum of $500-$600',  
    ↪'Sum of $600-$700',  
                'Sum of $700-$800', 'Sum of $800-$900', 'Sum of $900-$1000',  
    ↪'Sum of More than $1000']  
  
sum_total = 'Sum of Total number of delinquent residential accounts'  
  
demographics = ['pct_black', 'pct_hisp', 'pct_asian', 'pct_povt',  
    ↪'pct_immigrants', '% Renter Pop']
```

```

for pct, debt in zip(pct_debt_buckets, debt_buckets):
    acs[pct] = acs[debt] / acs[sum_total]*100

for dem in demographics:
    acs[dem] = acs[dem]*100

```

To simplify my code, I divided up much of the repetitive calculations into loops. My standardized buckets were listed together and looped, and I did the same for my demographic factors to multiply them by 100 for standardized percents in my data.

```
[10]: acs.head()
```

```

[10]:  Zip Codes  Sum of Less than $100  Sum of $100-$200  Sum of $200-$300  \
0      90001          5726.0          4937.0          1582.0
1      90002          3130.0          2152.0          1052.0
2      90003          1829.0          1880.0          1562.0
3      90004          2199.0          1659.0          1119.0
4      90005          1712.0          1107.0           598.0

      Sum of $300-$400  Sum of $400-$500  Sum of $500-$600  Sum of $600-$700  \
0              684.0          395.0          283.0          220.0
1              708.0          493.0          385.0          343.0
2             1169.0          941.0          782.0          673.0
3              735.0          502.0          387.0          279.0
4              401.0          281.0          175.0          146.0

      Sum of $700-$800  Sum of $800-$900  ...  Percent $100-$200  \
0              137.0          132.0  ...          33.129781
1              281.0          231.0  ...          20.009298
2              513.0          482.0  ...          13.753749
3              223.0          206.0  ...          19.315403
4               95.0           95.0  ...          21.731449

      Percent $200-$300  Percent $300-$400  Percent $400-$500  Percent $500-$600  \
0          10.616025          4.589988          2.650651          1.899074
1           9.781497          6.582985          4.583914          3.579730
2          11.427317          8.552198          6.884191          5.720974
3          13.028292          8.557457          5.844685          4.505763
4          11.739301          7.872006          5.516294          3.435414

      Percent $600-$700  Percent $700-$800  Percent $800-$900  \
0           1.476312          0.919340          0.885787
1           3.189214          2.612738          2.147838
2           4.923550          3.753018          3.526227
3           3.248341          2.596344          2.398417
4           2.866117          1.864939          1.864939

```

	Percent \$900-\$1000	Percent More than \$1000
0	0.650919	4.757751
1	1.868898	16.541144
2	2.933645	25.144488
3	1.909419	12.993364
4	1.138594	8.362780

[5 rows x 33 columns]

Above you'll see a quick check of my data. Next, I only wanted to keep columns of interest.

```
[11]: columns_drop = ['Sum of Less than $100',
    'Sum of $100-$200',
    'Sum of $200-$300',
    'Sum of $300-$400',
    'Sum of $400-$500',
    'Sum of $500-$600',
    'Sum of $600-$700',
    'Sum of $700-$800',
    'Sum of $800-$900',
    'Sum of $900-$1000',
    'Sum of More than $1000',
    'Sum of Total number of delinquent residential accounts']

acs = acs.drop(columns_drop, axis = 1)
```

The code I ran was a little shorter than listing the columns in their entirety. Then I checked my work.

```
[12]: acs.head()
```

```
[12]: Zip Codes      pop      mhhi  pct_black  pct_hisp  pct_asian  pct_povt  \
0      90001  58975.0  38521.0   8.864773  90.014413   0.218737  28.752380
1      90002  53111.0  35410.0  19.495020  78.463972   0.606277  32.860252
2      90003  72741.0  37226.0  22.182813  77.054206   0.353308  30.659739
3      90004  61586.0  48754.0   3.885623  51.113890  25.143701  18.060146
4      90005  39479.0  35149.0   6.137440  49.233770  35.079409  28.059292

      pct_immigrants  % Renter Pop  Percent Delinquent  ...  Percent $100-$200  \
0      40.763035      64.725731      25.268334  ...      33.129781
1      35.073714      62.199921      20.250042  ...      20.009298
2      37.544164      69.204438      18.791328  ...      13.753749
3      49.543727      80.390024      13.946351  ...      19.315403
4      59.135236      91.276375      12.903062  ...      21.731449

      Percent $200-$300  Percent $300-$400  Percent $400-$500  Percent $500-$600  \
0      10.616025      4.589988      2.650651      1.899074
```


1	9.781497	6.582985	4.583914	3.579730
2	11.427317	8.552198	6.884191	5.720974
3	13.028292	8.557457	5.844685	4.505763
4	11.739301	7.872006	5.516294	3.435414

	Percent \$600-\$700	Percent \$700-\$800	Percent \$800-\$900	\
0	1.476312	0.919340	0.885787	
1	3.189214	2.612738	2.147838	
2	4.923550	3.753018	3.526227	
3	3.248341	2.596344	2.398417	
4	2.866117	1.864939	1.864939	

	Percent \$900-\$1000	Percent More than \$1000
0	0.650919	4.757751
1	1.868898	16.541144
2	2.933645	25.144488
3	1.909419	12.993364
4	1.138594	8.362780

[5 rows x 21 columns]

0.1.5 Summary Statistics

Now that my data is cleaned and normalized, I want to get a better sense of how it looks. I'll calculate some summary stats.

```
[63]: Columns = ['pop',
                 'mhhi',
                 'pct_black',
                 'pct_hisp',
                 'pct_asian',
                 'pct_povt',
                 'pct_immigrants',
                 '% Renter Pop',
                 'Percent Delinquent',
                 'Percent Less than $100',
                 'Percent $100-$200',
                 'Percent $200-$300',
                 'Percent $300-$400',
                 'Percent $400-$500',
                 'Percent $500-$600',
                 'Percent $600-$700',
                 'Percent $700-$800',
                 'Percent $800-$900',
                 'Percent $900-$1000',
                 'Percent More than $1000']
```

```
acs[Columns].describe()
```

```
[63]:
```

	pop	mhhi	pct_black	pct_hisp	pct_asian	\
count	376.000000	376.000000	376.000000	376.000000	376.000000	
mean	28184.731383	82000.944149	4.188481	28.370609	13.309808	
std	21148.586559	36947.501222	5.912946	22.156196	14.473204	
min	202.000000	0.000000	0.000000	0.000000	0.000000	
25%	11313.500000	55786.750000	0.814591	11.902809	2.674377	
50%	26088.000000	74415.000000	2.005617	20.297812	8.145666	
75%	39401.000000	100338.500000	4.819766	40.371188	18.738314	
max	99284.000000	226450.000000	43.836772	97.971152	67.856784	

	pct_povt	pct_immigrants	% Renter Pop	Percent Delinquent	\
count	376.000000	376.000000	376.000000	366.000000	
mean	13.053572	22.855914	40.826504	0.372288	
std	8.330521	13.311688	18.082471	0.298633	
min	0.000000	0.000000	0.000000	0.000000	
25%	7.162339	12.851028	27.183280	0.095965	
50%	11.279916	21.250227	38.228502	0.322388	
75%	17.379622	31.288143	53.074624	0.613493	
max	66.666667	64.147483	92.803598	0.990099	

	Percent Less than \$100	Percent \$100-\$200	Percent \$200-\$300	\
count	314.000000	329.000000	327.000000	
mean	21.994885	23.360832	17.070812	
std	24.347467	21.544359	18.475945	
min	0.000000	0.000000	0.000000	
25%	0.000000	8.333333	7.310222	
50%	16.397849	19.444444	13.636364	
75%	32.948537	33.245383	19.667713	
max	100.000000	100.000000	100.000000	

	Percent \$300-\$400	Percent \$400-\$500	Percent \$500-\$600	\
count	319.000000	315.000000	303.000000	
mean	13.055340	57.238690	5.937663	
std	15.218510	844.713460	6.863027	
min	0.000000	0.000000	0.000000	
25%	5.509189	2.245201	1.316017	
50%	9.615385	6.060606	4.382470	
75%	15.849282	11.817279	8.021382	
max	100.000000	15000.000000	54.545455	

	Percent \$600-\$700	Percent \$700-\$800	Percent \$800-\$900	\
count	300.000000	298.000000	291.000000	
mean	5.066759	3.559053	3.232929	
std	8.567253	7.339987	9.000086	

min	0.000000	0.000000	0.000000
25%	0.573367	0.000000	0.000000
50%	3.041869	1.914782	1.665622
75%	6.107955	4.469320	3.341308
max	100.000000	100.000000	100.000000

	Percent \$900-\$1000	Percent More than \$1000
count	288.000000	313.000000
mean	2.596548	15.206207
std	8.837908	23.972368
min	0.000000	0.000000
25%	0.000000	1.136364
50%	1.015464	5.555556
75%	2.376241	18.340611
max	100.000000	100.000000

This is a lot of information in one place, but it can help me benchmark averages for debt and demographics across the state. Thus, in my further analysis of mapping demographics, I have this table as a baseline to see which Zip Codes are above, below the mean and median, and which Zips are towards the lower and upper bounds of the distribution. These stats will help guide my mapping as I determine Zip codes that have demographic factors that might put them at a higher risk of increased delinquent bill debt. Below I plot some of these relationships to better visualize them.

0.1.6 Incorporation of Demographics

I tried to use a for loop for these, but plotly.express did not produce the interactive maps within the function. The code ran, but none of the plots printed into the notebook, for reference, this is what I used:

```
def dem_scatter(dem): px.scatter(acs, x = dem, y = 'Percent Delinquent')

Demographics = ['mghi', 'pct_black', 'pct_hisp', 'pct_asian', 'pct_povt', 'pct_immigrants', '%
Renter Pop',]

for dem in Demographics: dem_scatter(dem)
```

Below you can see where I ran the plots by hand.

```
[24]: acs.drop(acs[acs['Percent Delinquent'] > 1].index, inplace = True)

px.scatter(acs,
            x = 'mghi',
            y = 'Percent Delinquent')
```

There is a very high concentration of zip codes with some degree of water bill debt between 30,000 and 100,000. This is a clear trend in the data, wherein lower income zip codes tend to have more water bill debt than higher income zip codes.

Now let's see if there are trends in the other demographic factors.

```
[25]: px.scatter(acs,
            x = 'pct_povt',
            y = 'Percent Delinquent')
```

Looking at the percentage of poverty, there is not necessarily a clear trend. There are definitely a high concentration of delinquent households in zip codes with poverty levels between 10-20%. However, when we dive deeper into this data, there might be a clearer trend.

The poverty rate in the US is around 12%, so if we consider anything higher than 12% as a high poverty area, there is some indication that there are a high concentration of zip codes with some degree of delinquency on their bills in moderately-high poverty rate zip codes. Of course, there needs to be a bit more examination of the data before we make a clearer judgement call.

```
[26]: px.scatter(acs,
            x = 'pct_black',
            y = 'Percent Delinquent')
```

There is a very slight trend here, but not a very obvious one. There are some zip codes with a higher percentage of Black populations that have higher bill debt (as you can see on the right hand side of the scatterplot). However, more analysis is needed.

```
[27]: px.scatter(acs,
            x = 'pct_hisp',
            y = 'Percent Delinquent')
```

It definitely seems like there are many more datapoints with a larger hispanic population than black population. It also does appear that populations with a greater percentage of hispanic populations also have fairly high levels of water bill debt delinquency (when compared to black populations) I think more analysis should be done, but there seems to be a clearer trend in this data than with the % black population.

```
[28]: px.scatter(acs,
            x = 'pct_asian',
            y = 'Percent Delinquent')
```

With the percent asian population, there is a less clear trend in the data, with a cluster towards the left hand side of the plot. If anything, the percent black, hispanic, and asian plots serve as clear comparisons and contrasts to the different ways these groups experience bill debt across the state.

```
[29]: px.scatter(acs,
            x = 'pct_immigrants',
            y = 'Percent Delinquent')
```

There is a fairly clear trend in this scatterplot, and I will incorporate immigrant populations into my map in the future.

```
[30]: px.scatter(acs,
            x = '% Renter Pop',
            y = 'Percent Delinquent')
```

There also appears to be a trend with renter population, and I will incorporate this into my mapping analysis as well.